

Revenue Model, Unit Economics, Financial Model & Pricing Test

Student	Jane Dias
Date	June 2026
Venture	AI Styling Feedback Platform

[DATA] = from published pricing, interview quote, or competitive research [ASSUMED] = reasoned from evidence; not directly tested

1. Revenue Model Selection

Models Eliminated

One-time purchase. The product builds a profile of proportions and style over time. A one-time payment for something that compounds with repeated use is structurally incoherent. Eliminated immediately. [ASSUMED]

B2B partnerships and affiliate / transaction fees. Allison said: "I understand the sales associate is just trying to make a sale... it's more reassuring when you ask a friend or family member." [DATA: Allison interview] Any model that pays me when she buys something collapses the trust the product is built on.

Per-use credits. My first recommendation after analyzing the interviews. Usage is episodic and decision-triggered, so per-session pricing seemed like a natural fit. The problem is what a shopping trip actually looks like: five items per fitting room, multiple stores, fifteen photos in one afternoon. Credits create scarcity anxiety at exactly the moment she needs to be fully in the decision. [ASSUMED: inference from interview behavior]

Advertising-supported freemium (Gate 1). At 500 monthly active users, ad revenue comes in at \$3–7.50/month [ASSUMED: CPM benchmarks]. That is not a business. The deeper problem: ads annoying enough to drive upgrades also undermine the trust that makes the advice worth paying for. I cannot optimize both simultaneously at early stage.

Free-with-no-memory gate (Gate 2). I proposed a stateless free tier, honest feedback but no memory, so the upgrade trigger would be organic. Rejected because removing memory defeats the product. The entire value is that it learns the user over time. A stateless tier is just another generic AI styling tool, which is exactly what Style DNA reviewers complained about. [DATA: Week 3 competitive research]

Model Selected

Free trial → monthly subscription at \$7/month. No ads at any tier.

Eight sessions free, full functionality, full profile building, full honest feedback. [ASSUMED: eight is a design decision; the right threshold is when the profile has enough data to be visibly valuable.] At the paywall, one decision: pay \$7/month or stop. One revenue party (the user),

one payment (monthly recurring via Stripe), one trigger (end of free trial). No advertiser relationship required. No transaction dependency.

The \$7/month price is not validated. I constructed it from three inputs: a direct variable cost floor of approximately \$3–4/month per active paying user [ASSUMED: derived from published LLM and infrastructure pricing]; a competitive range of \$7.99–\$12.99/month from Whering and Indyx [DATA: Week 3 competitive research]; and a value anchor from the interviews, every interviewee absorbs the cost of poor outfit decisions in wasted purchases rather than in money paid for solutions. [DATA: all five interviews] The Van Westendorp questions have not been run. The price is the most important open question in the model.

Why this beats the ad-supported model: The ad model required two businesses to work at the same time, an ad business that needs mass scale, and a subscription business with a structurally weakened conversion trigger. The subscription-only model closes at fifty paying subscribers. No scale dependency.

How Money Moves

- During the trial, no money moves toward me. Every session costs LLM inference before I see a dollar. At \$0.03/photo [ASSUMED], five photos per session, eight sessions, the trial cost is \$1.20 per user who completes it. At 10% conversion, ten trial users cycle through per paying subscriber, \$12 in acquisition LLM before the first payment. [ASSUMED]
- At the paywall, she pays \$7. Stripe takes 2.9% + \$0.30 [DATA: Stripe published pricing]. Stripe Tax adds 0.5% [DATA]. A 2% refund and chargeback rate [ASSUMED] reduces effective net revenue to approximately \$6.33 per subscriber per month.
- After subscription begins, the charge recurs monthly until cancellation. If she does not pay at the paywall, access stops. No fallback tier.

Evidence For

- The behavior already exists. All five interviewees use photos as part of outfit decisions, Mya, Danaï, and Allison all self-photograph before committing to a purchase. The trial converts an existing behavior into a product interaction. [DATA: all five interviews]
- The model closes at small scale. Fifty paying subscribers at \$7/month generates \$317 in net revenue, nearly enough to cover \$343 in non-labor fixed costs. No advertising scale required. [ASSUMED: math from unit economics model]
- The competitive range is real. Whering is \$7.99/month. Indyx is \$9–\$12.99/month. [DATA: Week 3] \$7 sits at the low end.
- The value anchor survives the budget constraint. Mya bought a Target tube top she never wore. Aly accumulated non-returnable sale items. [DATA: interviews] The right frame is not "is \$7 worth it for an app" but "is this worth more than one purchase you will regret."

Evidence Against

- No interviewee has been asked what they will pay. The Week 4 assignment said it directly: "there is limited information about whether or not the target customer will actually pay for this product." [DATA: Week 4] The price is constructed, not observed.
- Every current substitute is free. Allison texts family. Mya FaceTimes her mom. Danai sends photos to her husband. [DATA: all five interviews] I am asking budget-limited women to pay monthly for something they currently get at \$0.
- Usage is episodic. Subscriptions reward continuous use. In months where she is not shopping, she is paying for a subscription she did not use. That is the definition of the subscriptions people cancel. [ASSUMED]
- The competitive set failed in this category. Style DNA reviewers said the advice was not worth the cost. Alle App raised \$3M and failed for lack of product-market fit. [DATA: Week 3 competitive research]

2. Unit Economics

Detailed figures are in the Excel model (unit-economics.md and the Excel Unit Economics sheet). This section explains the reasoning.

Revenue Per Customer

\$7 gross per subscriber per month. After Stripe processing (2.9% + \$0.30) [DATA], Stripe Tax (0.5%) [DATA], and a 2% refund rate [ASSUMED], approximately **\$6.33 arrives per subscriber per month**. Trial users generate zero revenue.

Costs and Why They Are What They Are

LLM inference is the largest and least predictable direct cost. Every photo costs approximately \$0.03 [ASSUMED: Claude Sonnet pricing estimate]. Paying subscribers generate ~43 photos/month [ASSUMED] = \$1.29. Trial users generate 5 photos per session across 8 sessions = \$1.20 per trial user. The critical number: at 10% conversion and 16.7% churn, 1.67 trial users cycle through per paying subscriber retained, allocating \$2.00 in trial LLM cost per paying subscriber before they generate a dollar of profit.

Variable infrastructure scales with all active users. At \$0.25/user/month [ASSUMED: marginal AWS cost], a combined population of ~260 users in month twelve costs approximately \$65 that month. The pessimistic scenario assumes \$0.35/user due to higher server demands.

Refunds and chargebacks at 2% [ASSUMED: industry midpoint] reduce net revenue by approximately \$10/month at month-twelve scale. Stripe processes refunds automatically with no developer recourse. [DATA: Stripe documentation] It is in the model because it affects net revenue, not because it is large.

Stripe Tax at 0.5% [DATA] automates sales tax collection across ~30 US states that require it for SaaS subscriptions. Not optional.

Contribution Margin Per Subscriber

\$6.33 revenue – \$1.54 variable costs – \$2.00 allocated trial LLM = \$2.79 per subscriber per month. That is what each additional paying subscriber contributes toward covering fixed costs. At \$1,378/month in fixed costs, I need 496 paying subscribers to break even.

My Time and Why It Changes the Model Materially

Founder labor covers all founder time at \$15/hour [DATA: per course instruction].

Base scenario: 9 hours/week of product, support, ops, and admin [ASSUMED], plus 30 hours/month of TikTok and Instagram content creation [ASSUMED: derived from unit-economics.md]. Combined: \$1,035/month. This is the single largest fixed cost line.

Content creation is not a separate expense. It is founder time, the same \$15/hour as everything else. If I stop posting, trials stop, and the entire model collapses. I cannot account for the output without paying for the input.

Total founder labor in the base scenario: \$1,035/month.

Fixed Costs That Run Regardless

Beyond founder labor: server infrastructure \$70/month [ASSUMED: AWS t3.medium + RDS]; legal and compliance \$75/month [ASSUMED: privacy policy, TOS, GDPR/CCPA amortized]; cyber liability, general liability, and E&O insurance \$165/month [ASSUMED: midpoint of quoted range]; software tools + LLC fees \$33/month [ASSUMED]. Total non-labor fixed costs: \$343/month. Total with founder labor and content creation: \$1,378/month.

What the Cost Audit Revealed

The AI's first pass excluded content creation entirely, classifying it as optional CAC spend that could be paused and assigned \$0. After correction, content creation was included as 30 hours/month of founder time at \$15/hr, part of the combined founder labor line. Fixed costs in the base scenario are \$1,378/month. The full audit trail is in the Excel.

The lesson: AI will zero out any cost that seems discretionary. Content creation is not discretionary, it is the mechanism that generates trial volume. If I stop posting, trials stop. If trials stop, the entire model collapses.

Additional findings from the audit: cyber liability insurance is non-optional (I am storing personal photos of users in fitting rooms, a breach is materially different liability than a breach of email addresses); App Store distribution is excluded because Stripe direct avoids Apple's 30% first-year fee [DATA: Apple published fee structure]; taxes are not in the projections (self-employment tax at 15.3% [DATA: IRS] would apply immediately to any profitable scenario); content moderation adds approximately \$0.012/user/month [ASSUMED], negligible at year-one volume but a real legal exposure from omitting the capability.

3. Financial Model Summary

Month-by-month detail for all three scenarios is in the Excel model (tabs: Year 1 Financial Model, Pessimistic, Optimistic).

	Pessimistic	Base	Optimistic
Price / month	\$7	\$7	\$9
Trial signups, month 12	95 [ASSUMED]	190 [ASSUMED]	475 [ASSUMED]
Sell-through (conversion)	5% [ASSUMED]	10% [ASSUMED]	20% [ASSUMED]
Monthly churn	25% [ASSUMED]	16.7% [ASSUMED]	10% [ASSUMED]
Items (paying subs, M12)	13	73	462
Storage (server / month)	\$90 [ASSUMED]	\$70 [ASSUMED]	\$70 [ASSUMED]
Year 1 net revenue	\$453	\$2,589	\$19,295
Year 1 variable costs	(\$770)	(\$2,185)	(\$7,523)
Year 1 fixed costs (incl. content creation)	(\$19,896)	(\$16,536)	(\$14,976)
Net, at \$15/hr [ASSUMED]	(\$20,214)	(\$16,132)	(\$3,205)
Net, at \$0/hr (content=\$0, labor=\$0) [ASSUMED]	(\$4,674)	(\$3,712)	\$7,655

What This Tells Me

All three scenarios lose money at \$15/hr in year one.

- The pessimistic scenario cannot break even at any scale. At 5% conversion and 25% churn, five trial users cycle through per paying subscriber. With full pessimistic parameters (7 photos per trial session, 5% refund rate, \$0.35 infrastructure): net revenue is \$6.11, trial LLM is \$8.40 per subscriber (\$1.68 per trial user × 5), and direct variable costs are \$1.85. Contribution margin is -\$4.14 per subscriber. Every new subscriber deepens the loss. This is not a cost problem, it is a conversion and churn problem.
- Storage (server infrastructure) differs by scenario, \$90/month pessimistic, \$70/month base and optimistic, but neither figure moves profitability materially. The real lever is conversion rate.
- The path to breakeven at base behavioral assumptions: 496 paying subscribers, sustained by ~828 trial signups per month. My model peaks at 190 trials/month in month 12. Breakeven is a year-two question at base assumptions. At optimistic behavioral assumptions (\$9 price, 20% conversion, 10% churn), breakeven is 205 subscribers, reachable around month nine of year one.

What Would Need to Change

- Price is the most powerful single lever, but not at base trial volume. No price change reaches breakeven in year one at 190 trials/month peak. The optimistic scenario (\$9, better behavioral assumptions) reaches breakeven around month nine.
- **Conversion and churn** determine structural viability. The pessimistic scenario fails regardless of price or scale. Each 1-point improvement in monthly churn (16.7% → 15.7%)

reduces the breakeven subscriber count by approximately 30. These are product design questions.

4. Pricing Test Plan

Hypothesis

Women aged 22–35 who regularly seek second opinions on outfit choices will pay \$9/month for AI styling feedback, because they have no reliable, bias-free feedback mechanism, and \$9 is less than what they spend on items they return. [ASSUMED: every element is what this test is designed to check]

I am testing \$9, not \$7. At \$9 with optimistic behavioral assumptions (20% conversion, 10% churn), breakeven requires 205 paying subscribers, a level the optimistic projection reaches mid-year. At \$7 with base behavioral assumptions, breakeven requires 496 subscribers and ~828 trial signups per month, a materially longer path. The competitive range runs \$7.99–\$12.99/month [DATA: Whering and Indyx, Week 3]. If Mya and Aly will not pay \$9 but will pay \$7, that tells me my price ceiling is below what the optimistic model needs.

Test Method: Concept Card + Van Westendorp Interview

A Wizard of Oz test requires real payment and real delivery, more infrastructure than I need right now. A direct ask ("would you pay for this?") produces false positives; people say yes to be polite. This test shows a concrete concept card to each interviewee, then runs the Van Westendorp four-question sequence. Van Westendorp does not ask "would you pay?", it asks for four price thresholds that together reveal the acceptable price range. No payment is collected. The data is where each person's thresholds fall relative to \$9.

The Four Questions

Ask these in order. Do not say "\$9" before asking, any number I introduce first becomes an anchor.

- (1) "At what price would this feel so cheap that you'd question whether it actually works?" (finds the credibility floor)**
- (2) "At what price would this feel like a genuinely good deal?" (finds the sweet spot)**
- (3) "At what price would it start to feel expensive, but you'd still consider it?" (finds the stretch zone)**
- (4) "At what price would it be too expensive, no matter how good it is?" (finds the ceiling)**

After all four answers: "After 8 free sessions where the app has learned your proportions and what works on you, would you pay \$9 to keep that profile?" Write down their exact words.

The Concept Card

Print or show on phone. One card per interview.

AI Styling Feedback

Honest feedback on any outfit, fitting room, mirror at home, before you go out.

Not "you look great." Specific: does this work on your frame? Would you actually wear it again?

How it works

- Take a photo. Get honest feedback in under a minute.
- The more you use it, the more it learns your proportions and what works for you.
- No sales associate. No friend who is being polite. No agenda.

Free trial: 8 sessions, no card required, no commitment

After that: \$9/month to keep your profile and everything it has learned about you

\$9/month is less than most items you have bought and never worn.

Steps

Monday, June 23: Prepare: Print or save the concept card. Write down each person's name, their clothing budget from the interviews, and their primary pain point. I am not going in cold, I know these people.

Tuesday, June 24: Interview Mya and Aly: My primary targets. Both named the honesty gap explicitly [DATA: interviews]. Both have limited clothing budgets [DATA]. Show the card. Let them read it. Then ask the four Van Westendorp questions in order. Record their exact numbers. Only after all four, ask the \$9 follow-up.

Wednesday, June 25: Interview Allison and two others: Allison is a secondary target, older, larger budget, partially self-solved [DATA: interview]. Her answers tell me whether the product skews toward a different segment. The other two should be women I have not interviewed before who match the target profile and would tell me no if they meant no.

Thursday, June 26: Follow up on unclear answers: If anyone gave a range instead of a number, go back and ask which end they would act on. One follow-up per person maximum.

Friday, June 27: Tally: For each of the five: what they said, their reaction to \$9, their Van Westendorp range. Three sentences on what I learned.

Success and Failure Criteria

Signal	What It Looks Like
Yes	3 or more of 5 give a Van Westendorp "too expensive" threshold above \$9. When asked the follow-up, "would you pay \$9 after 8 free sessions?", they say yes without being prompted to reconsider. At least one asks a follow-up question about the product.
Weak yes	2 of 5 say yes to \$9. The others' Van Westendorp

	thresholds cluster between \$7 and \$9, they find value but not at \$9. The price is at the ceiling. Test \$7 with the same group before committing.
Price no	0–1 say yes to \$9, but Van Westendorp thresholds from the others come in at \$5–8. They find value but the ceiling is below what the model needs. Run Plan B. This is a price problem, not a product problem.
Product no	0–1 say yes to \$9 and thresholds are at \$0–4. Ask directly: "What would make this worth paying for?" If they cannot answer, the barrier is product-market fit, not price.

Success: 3 of 5 interviewees place \$9 within their acceptable range (above their credibility floor, below their ceiling), and at least 2 of 5 say yes to the direct \$9 follow-up question.

Failure: fewer than 3 place \$9 in range, or no one says yes to the follow-up.

If the Test Fails

1. Test \$7. If Van Westendorp thresholds cluster between \$7 and \$9, return to the same five people with the card repriced at \$7. If 3+ say yes, I have a price ceiling, the model can work at \$7, but breakeven rises from ~205 subscribers (optimistic, \$9) to ~496 (base behavioral, \$7), and required trial volume rises from ~103 to ~828/month. That is a materially longer path. Name it explicitly.

2. Test a shorter free trial. If people say \$9 is fine but they would not trust the advice until they had used it more, the friction is not price, it is trial length. Ask: "How many sessions would you need before you'd feel confident paying?" If the answer is 3 or 4, trial LLM costs drop and the conversion moment arrives sooner.

3. Test subscription frame vs. per-use. If people resist \$9/month as a commitment but respond positively to "10 sessions for \$9," the barrier is recurring commitment, not the dollar amount. A credit bundle is operationally harder but may unlock episodic shoppers who are averse to subscriptions they might not use every month.

Alternative Test: Concierge

The concept card + Van Westendorp interview above tests stated preference, what people say they would pay. A concierge test tests revealed preference, whether they actually use the product and pay when the moment arrives. Run Van Westendorp first to confirm the price is in range, concierge second to confirm people actually convert.

What It Is

I manually deliver the styling feedback myself, no app, no AI. I am the product. The pitch:
"Send me any outfit photo, fitting room, mirror, before you go out, and I'll give you honest feedback within a few hours. After 8 sessions I'll ask you to pay \$9 to keep going." No app. iMessage or Instagram DM. At session 8: "You've hit the free limit. To keep getting feedback, it's \$9, Venmo or CashApp." I record exactly what happens.

What to Log Each Session

What they sent: fitting room shot, mirror photo, event outfit, regret purchase, the category matters

What I said: my exact feedback, how long it took me to write it

Their reaction: did they reply, act on it, send another photo soon after

Paywall moment: what they said or did when asked to pay \$9 at session 8

What This Tells Me That Van Westendorp Does Not

- Whether they actually send photos, or whether friction stops them before they start.
- Whether the feedback changes their behavior, do they mention acting on it.
- How long it takes me to give one good response (15–20 minutes per photo is a real operational number that does not appear in my model).
- Which use case gets the most photos, fitting room indecision, event outfits, past purchases, which tells me where the product is strongest.
- Whether they pay at the paywall moment, not whether they said they would.

The Tradeoff

Van Westendorp takes 30 minutes per person. A concierge test takes 2–4 weeks and roughly 2–3 hours of my time per participant before I see a conversion decision. Run the interview test first. If the price thresholds come back above \$9, the concierge test is the logical next step.